

# CS190C Lec1

Word Embedding & Language Modeling

# Overview

- How to embed words
- What is language modeling
- Some naive and early language models
  - N-grams
  - RNN
  - LSTM
- Thinking

# **PART1: How to embed words?**

Problem: Computer cannot understand natural language.....

- We should try to convert them, such as words, into digital form.
- How to represent words formally?

## A naive idea

- Like a dictionary, each word has its position in it
- Can we represent a word using a certain number?
  - That is: each word has its "one-to-one mapping value"
- Just like this:

Word	I	love	Natural	Language	Processing	" "	.
Value	0	1	2	3	4	5	6

"I love Natural Language Processing."  $\Rightarrow$  0,5,1,5,2,5,3,5,4,6

# Mapping has its defect

- Can we have a method to infer this word's part of speech or meaning, just based on the mapping value?
- For example: If you only receive a string of numbers: `0,5,1,5,2,5,3,5,4,6`, can you successfully infer the meaning of this sentence?
- Mathematically, one-to-one mapping can be understood as: **The semantics of different words constitute a one-dimensional vector space!**
  - Each word is a 1-dim vector in this 1-dim space
  - Using 1-dim vector space to represent natural language is obviously not enough!

## Enlarge the dimension?

- If we use 2-dim space to represent some words?
  - x-axis represent "area", y-axis represent "population"
  - Can we use 2-dim vectors to approximately represent: China   India   Canada  
Luxembourg ? (Plot a graph of it)
- Similarly, the higher the dimension, the richer the semantics it can represent.
- For example, GPT-3 uses 2048-dim to represent words.

# What is word embedding?

So far we know that:

- We can use n-dim vector to represent the meaning of the word formally.
- Each dim can be understand as a kind of semantics, it just like one more dimension of space in linear algebra.

We call vectors of word **Word Embedding**.

- Encoding words with appropriate word embeddings is one of the key points in natural language processing technology
- In later lecs, we will introduce some important encoder models.

## PART2: What is language modeling?

# Look at a LLM

Qwen3-Max

NLP stands for **Natural Language Processing**. It is a subfield of **artificial intelligence (AI)** and **computational linguistics** that focuses on enabling computers to understand, interpret, generate, and interact with **human language** in a meaningful and useful way.

## Key Goals of NLP:

1. Understanding: Extract meaning from human language (e.g., sentiment analysis, intent recognition).
2. Generation: Produce human-like text (e.g., chatbots, automated reports).
3. Translation: Convert text or speech from one language to another (e.g., Google Translate).
4. Interaction: Enable seamless communication between humans and machines using natural language (e.g., voice assistants like Siri or Alexa).

## Common NLP Tasks:

- Tokenization: Splitting text into words, phrases, or symbols.
- Part-of-Speech Tagging: Identifying nouns, verbs, adjectives, etc.
- Named Entity Recognition (NER): Detecting entities like people, organizations, dates.
- Sentiment Analysis: Determining if a sentence expresses positive, negative, or neutral emotion.
- Machine Translation: Translating text between languages.
- Text Summarization: Condensing long documents into shorter versions.
- Question Answering: Systems that answer questions posed in natural language.
- Speech Recognition & Synthesis: Converting speech to text and vice versa.

What is NLP?

- Input a prompt (a string of words)
- Give an answer based on the prompt

How does it generate words, finally forms an answer?

- Like normal speaking of human, each word should generate after formal words, and based on it logically.
- That is: based on old words, and generate new words.
- This called Language Modeling!

## Tips: Difference between **Language Model** and **Language Modeling**

- Language Model: Is a tool to generate certain answer based on prompts.
- Language Modeling: Is the methods to generate "new words" based on "old words".

So, what's the "methods"?

## General methods

There are maybe several words suitable to be the generated new word.....

- But different word may have different levels of suitability.
- We can try to model "levels of suitability" into probabilistic distribution.
- That is, find a way to calculate the probability of generating a certain word as the "new word", based on "old words".

# Language Modeling Methods.....

These are different models, using different certain methods to calculate the probabilistic distribution.

- N-grams
- RNN, LSTM (An optimized architecture based on RNN)
- Transformer
- GPT, BERT .etc (Based on Transformer)

## PART3: N-grams

## Another naive idea

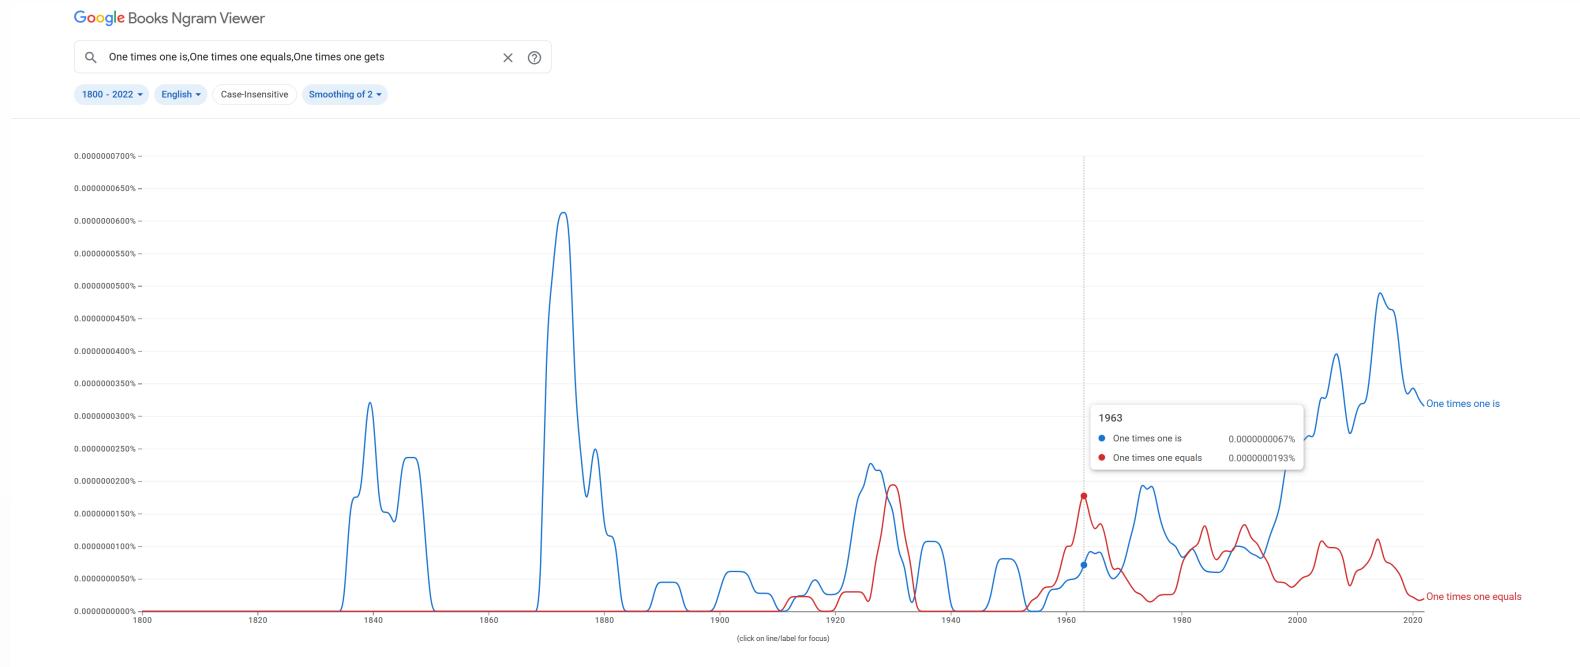
Simply consider: If we just focus on a fix window of old words to generate a new word, it can still work at most of time.

For example:

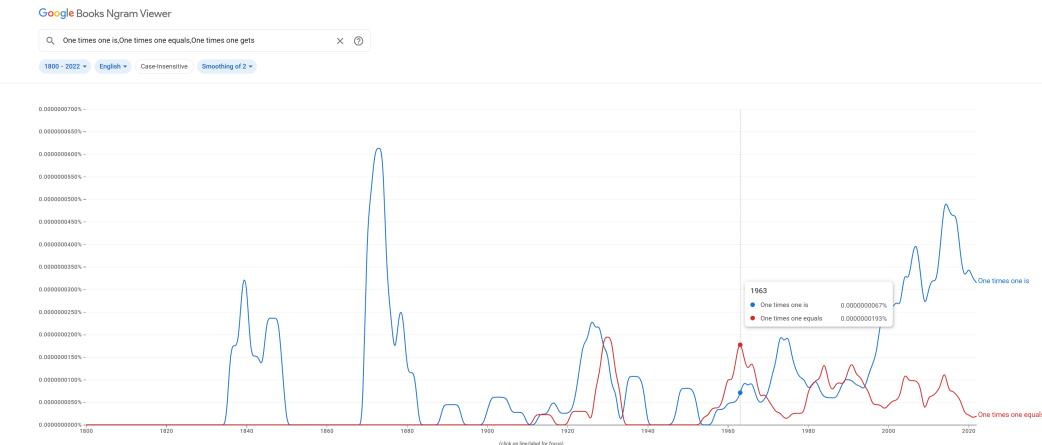
- We only focus on latest 3 old words to generate a new word
- We call it **4-grams**, which means a fixed window contain old words + a new word has 4 words.
- Prompt: "Let's calculate simple multiplication! One times one....."

# A naive 4-grams

- old words focused: One times one
- Generate new word directly use statistical laws!
- We can use <https://books.google.com/ngrams/>



# A naive 4-grams



- We can try all words in vocabulary
- We can also get  $P(\text{new}, \text{old})$  (marginal probability of 4-grams)
- What we want to model is  $P(\text{new}|\text{old}) \propto P(\text{new}, \text{old})$

- Suppose we've modeled  $P(w_i|\text{old})$
- We choose  $\text{argmax}_{w_i} P(w_i|\text{old})$  or sample words according to the distribution
- Suppose we decide `is` to be the new word
- Sentence now: One times one `is`
- Next turn: use `times one is` to generate a new word, and so on.

## Pros and cons?

- Actually it is quite simple, and do not need much calculation (especially model training)
- If you really use N-grams to generate text, what will happen?

# Generate text?

*today the price of gold per ton , while production of  
shoe lasts and shoe industry , the bank intervened just  
after it considered and rejected an imf demand to  
rebuild depleted european stocks , sept 30 end primary  
76 cts a share .*

Surprisingly grammatical!

...but **incoherent**. We need to consider  
more than three words at a time if we  
want to model language well.

But increasing  $n$  worsens sparsity  
problem, and increases model size...

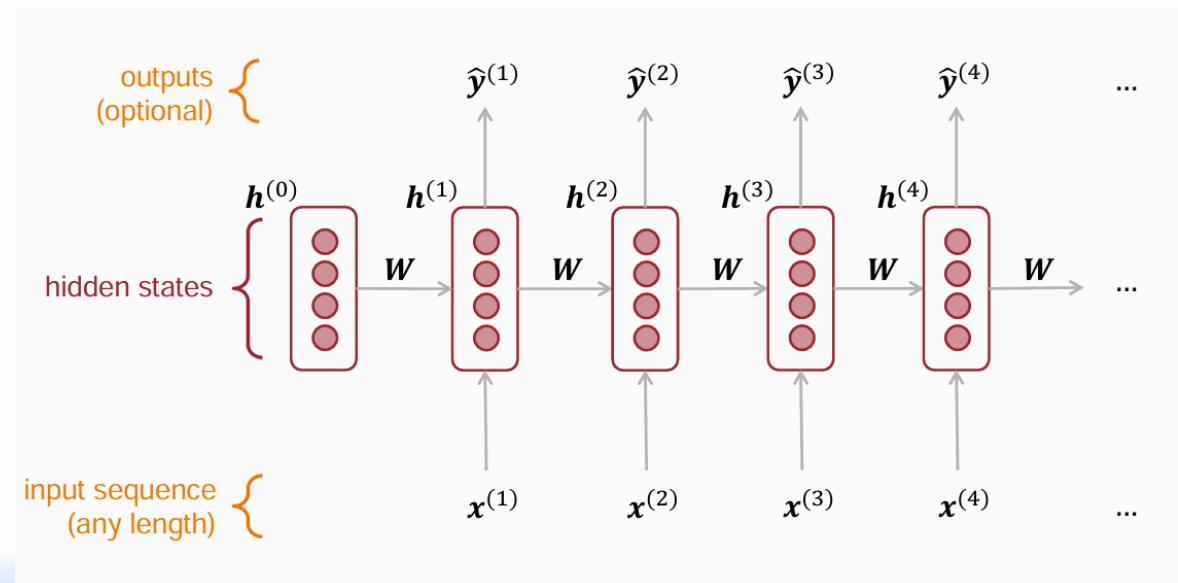
- For grammar.....
  - Subject-predicate/verb-object.....  
part-of-speech collocation follows  
statistical rules
- But for semantics.....
  - Missing of global context
  - One mistake will cause  
accumulation of subsequent  
errors

## PART4: Recurrent Nerual Network (RNN)

# How to make use of global context?

Encodding.....

- We encode words as vectors.....
- Can we encode text as vectors too?
- If we get the encoded vector of text, we can generate new words with "understanding" of text!



## How to encode words and text?

- For words.....
  - At first, each word still encoded with single number.
  - We hope to train a matrix  $W_e \in R^{d*|V|}$ , each column is the embedding of i-th word.
  - So for i-th word, we can get its embedding through  $W_e e$ ,  $e$  is the one-hot vector of this word (only get value 1 in i-th element)  $\Rightarrow W_e e$  is the i-th column of  $W_e$

## How to encode words and text?

- For text.....
  - Like reading word by word, with each word read, the understanding of the text will be more substantial.
  - When a new word read: the "understanding" of the text will mixed with: Former understanding of the text, and the information of the new word.
  - That is: the embedding of the text in this time step should mix former embedding of text and embedding of the new word together.

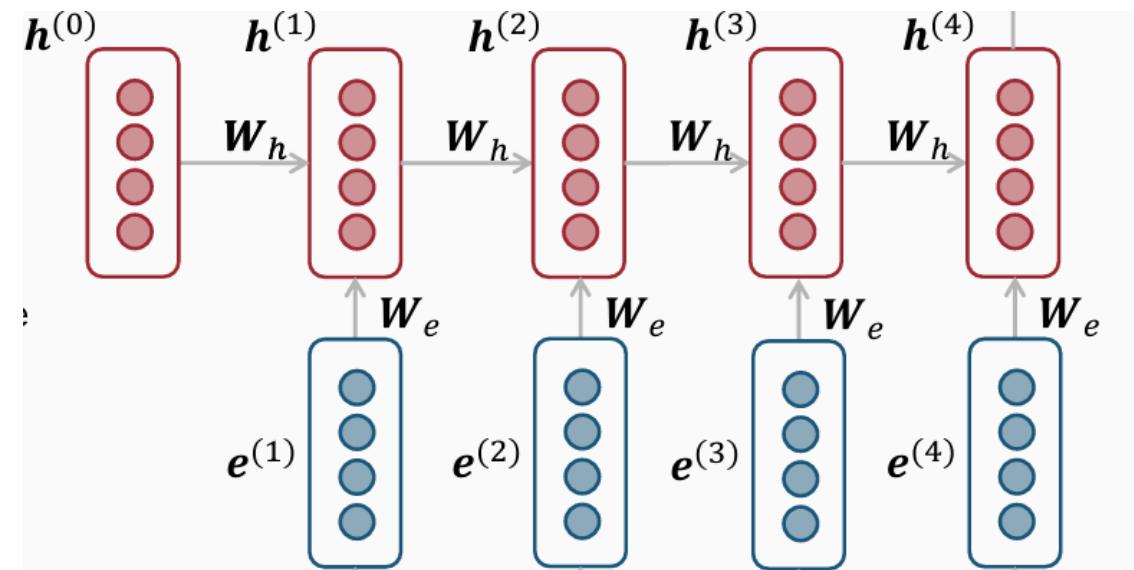
## How to encode words and text?

Let embedding of text in step-t be  $h^t$

- Inherit former embedding:  $W_h h^{t-1}$ . We hope to train  $W_h$ , which stands for how to proper inherit former information
- New word's information:  $W_e e^t$
- Combine:  $W_h h^{t-1} + W_e e^t + b_1$ . ( $b_1$  is the optional bias term)
- Add a nonlinear activation (usually use sigmoid)

## How to encode words and text?

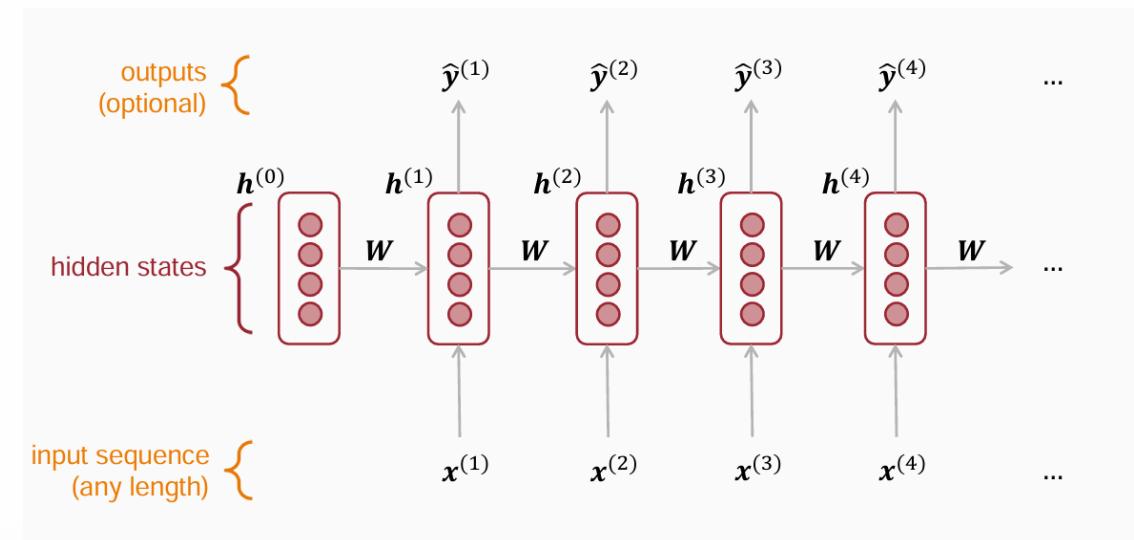
- $h^t = \sigma(W_h h^{t-1} + W_e e^t + b_1)$
- Parameters to train:
  - $W_e$
  - $W_h$
  - $b_1$



# How to language modeling?

Just use a linear activation to  $h^t$ , generate the probability distribution of the new words!

- $\hat{y}^t = \text{Softmax}(Uh^t + b_2)$
- Parameters to train:
  - $U$ : the linear activation matrix
  - $b_2$ : the optional bias term



# Implement a simple RNN

[https://github.com/kuangpenghao/NLP\\_models\\_by\\_hand/blob/main/toy\\_RNN.py](https://github.com/kuangpenghao/NLP_models_by_hand/blob/main/toy_RNN.py)

- 3 training sentences
- Train RNN
- Input sentence except the last word
- Hope to output the correct word

```
epoch:0200,loss:0.715562
epoch:0400,loss:0.200380
epoch:0600,loss:0.097146
epoch:0800,loss:0.060034
epoch:1000,loss:0.064022
input:['i', 'like'],output:dog
input:['i', 'love'],output:coffee
input:['i', 'hate'],output:milk
```

# Implement a simple RNN

- main function

```
if __name__=="__main__":
    config=TextRNNConfig()
    model=TextRNN(config)

    trainer=TextRNNTrainer(config,model)
    predictor=TextRNNPredictor(config,model)

    trainer.train()

    test_sentences=["i like dog", "i love coffee", "i hate milk"]
    for sentence in test_sentences:
        sentence=sentence.split()
        input_sentence=sentence[:-1]
        predicted_word=predictor.predict(input_sentence)
        print(f"input:{input_sentence},output:{predicted_word}")
```

# Implement a simple RNN

```
class TextRNNConfig:  
    def __init__(self):  
        self.n_hidden=5  
        self.sentences=["i like dog", "i love coffee", "i hate milk"]  
  
        word_list=' '.join(self.sentences).split()  
        word_list=list(set(word_list))  
        self.word_dict={w:i for i,w in enumerate(word_list)}  
        self.number_dict={i:w for i,w in enumerate(word_list)}  
  
        self.n_class=len(word_list)  
  
        self.batch_size=2  
        self.learning_rate=0.001  
        self.epoches=1000  
        self.interval=200
```

# Implement a simple RNN

```
class TextRNN(nn.Module):
    def __init__(self, config):
        super(TextRNN, self).__init__()
        self.config = config
        self.rnn = nn.RNN(self.config.n_class, self.config.n_hidden)
        self.W = nn.Linear(self.config.n_hidden, self.config.n_class, bias=True)

    def forward(self, X):
        X = X.transpose(0, 1)
        ori_hidden = torch.zeros(1, X.shape[1], self.config.n_hidden)
        outputs, last_hidden = self.rnn(X, ori_hidden)
        output = outputs[-1]
        output = self.W(output)
        return output
```

# Implement a simple RNN

```
class TextRNNDataset(Dataset):
    def __init__(self, config):
        super(TextRNNDataset, self).__init__()
        self.config=config

    def __len__(self):
        return len(self.config.sentences)

    def __getitem__(self, idx):
        sentence=self.config.sentences[idx]
        words=sentence.split()
        words=[self.config.word_dict[i] for i in words]

        input_idx=words[:-1]
        one_hot=np.eye(self.config.n_class)[input_idx]
        input_one_hot=torch.tensor(one_hot,dtype=torch.float32)

        target_idx=words[-1]
        target_idx=torch.tensor(target_idx,dtype=torch.int64)

        return input_one_hot,target_idx
```

# Implement a simple RNN

```
class TextRNNTrainer:  
    def __init__(self, config, model):  
        self.config=config  
        self.model=model  
        self.loss_function=nn.CrossEntropyLoss()  
        self.optimizer=torch.optim.SGD(model.parameters(), lr=self.config.learning_rate, momentum=0.9)  
  
        self.datagetter=TextRNNDataset(config)  
        self.dataloader=DataLoader(self.datagetter, batch_size=self.config.batch_size, shuffle=True)  
  
    def train(self):  
        for epoch in range(self.config.epochs):  
            for input_batch, target_batch in self.dataloader:  
                self.optimizer.zero_grad()  
                output=self.model(input_batch)  
                loss=self.loss_function(output, target_batch)  
                loss.backward()  
                self.optimizer.step()  
            if (epoch+1)%self.config.interval==0:  
                print(f"epoch:{epoch+1:04d}, loss:{loss.item():.6f}")
```

# Implement a simple RNN

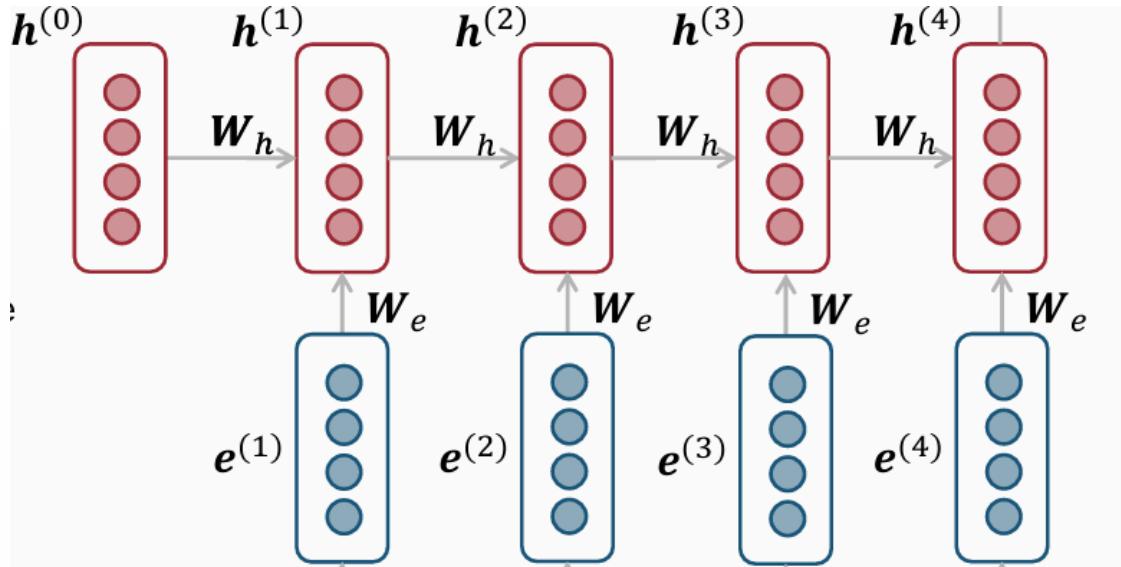
```
class TextRNNPredictor:  
    def __init__(self, config, model):  
        self.config=config  
        self.model=model  
  
    def predict(self, input_sentence):  
        with torch.no_grad():  
            word=[self.config.word_dict[i] for i in input_sentence]  
            word=np.eye(self.config.n_class)[word]  
            word=torch.tensor(word,dtype=torch.float32).unsqueeze(0)  
  
            output=self.model(word)  
            outcome=output.max(1,keepdim=False)[1]  
            predicted_word=self.config.number_dict[outcome.item()]  
  
        return predicted_word
```

## Pros and cons?

- Actually, it can make use of the information of the global text.....
- When the text is extremely long, step  $t$  is very large?

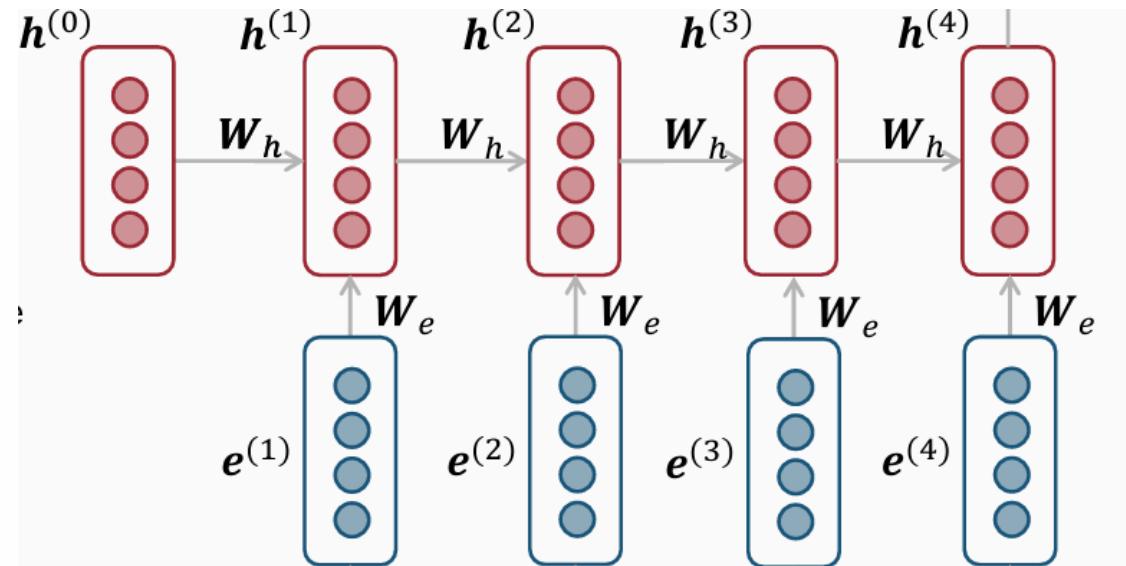
# Update parameters ( $W_h$ )

- $\frac{\partial L_T}{\partial W_h} = \sum_{t=1}^T \frac{\partial L_T}{\partial h^t} \frac{\partial h^t}{\partial W_h}$
- $\frac{\partial L_T}{\partial h^t} = \frac{\partial L_T}{\partial h^T} \frac{\partial h^T}{\partial h^{T-1}} \cdots \cdots \frac{\partial h^{t+1}}{\partial h^t}$
- $\frac{\partial h^{t+1}}{\partial h^t} = \frac{\partial \sigma(W_h h^t + W_e e^{t+1} + b_1)}{\partial h^t}$   
 $= \sigma'(W_h h^t + W_e e^{t+1} + b_1) W_h$
- We usually use sigmoid function as  $\sigma$ .  
 Then  $\sigma'(z) = \sigma(z)(1 - \sigma(z)) \in (0, 0.25]$
- $\frac{\partial L_T}{\partial h^t} = \frac{\partial L_T}{\partial h^T} \prod_{k=t+1}^T [\sigma'(z_k) W_h]$



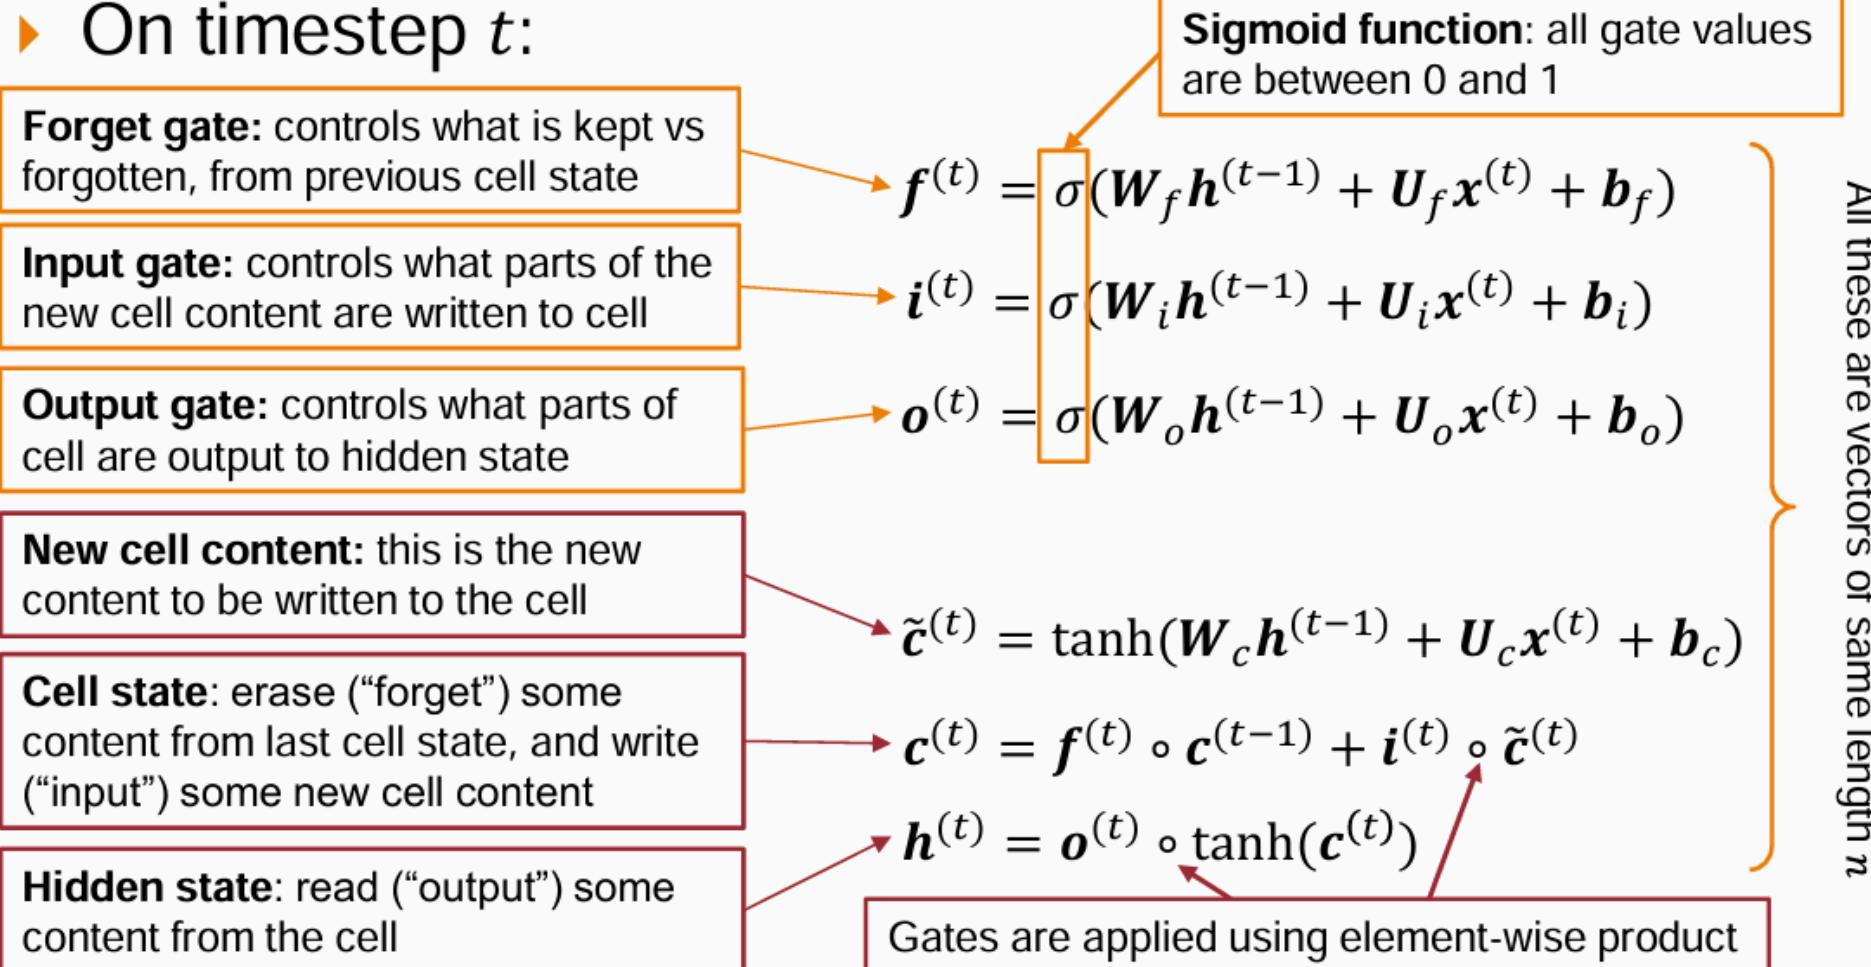
## Update parameters ( $W_h$ )

- $\frac{\partial L_T}{\partial h^t} = \frac{\partial L_T}{\partial h^T} \prod_{k=t+1}^T [\sigma'(z_k) W_h]$
- If  $\|W_h\| < 1$ , each term must  $< 1$   
 $\Rightarrow \frac{\partial L_T}{\partial h^t} \rightarrow 0$ , which may leads to  
 $\frac{\partial L_T}{\partial W_h} \rightarrow 0$
- Likewise, if  $\|W_h\|$  is large enough,  
 $\frac{\partial L_T}{\partial W_h} \rightarrow \infty$
- Called vanishing gradient & exploding gradient

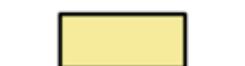
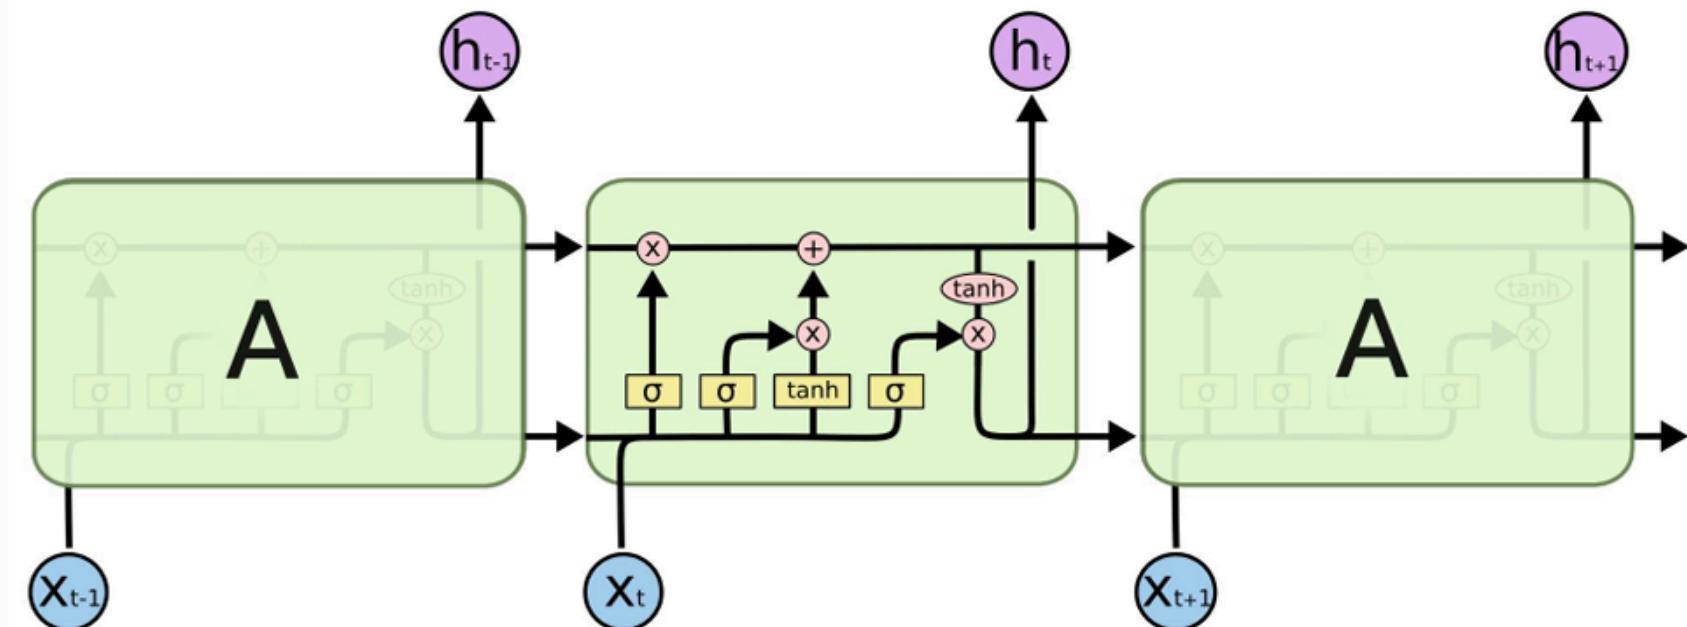


## PART5: Long Short Term Memory (LSTM)

- ▶ We have a sequence of inputs  $\mathbf{x}^{(t)}$ , and we will compute a sequence of hidden states  $\mathbf{h}^{(t)}$  and cell states  $\mathbf{c}^{(t)}$ .
- ▶ On timestep  $t$ :



► You can think of the LSTM equations visually like this:



Neural Network  
Layer



Pointwise  
Operation



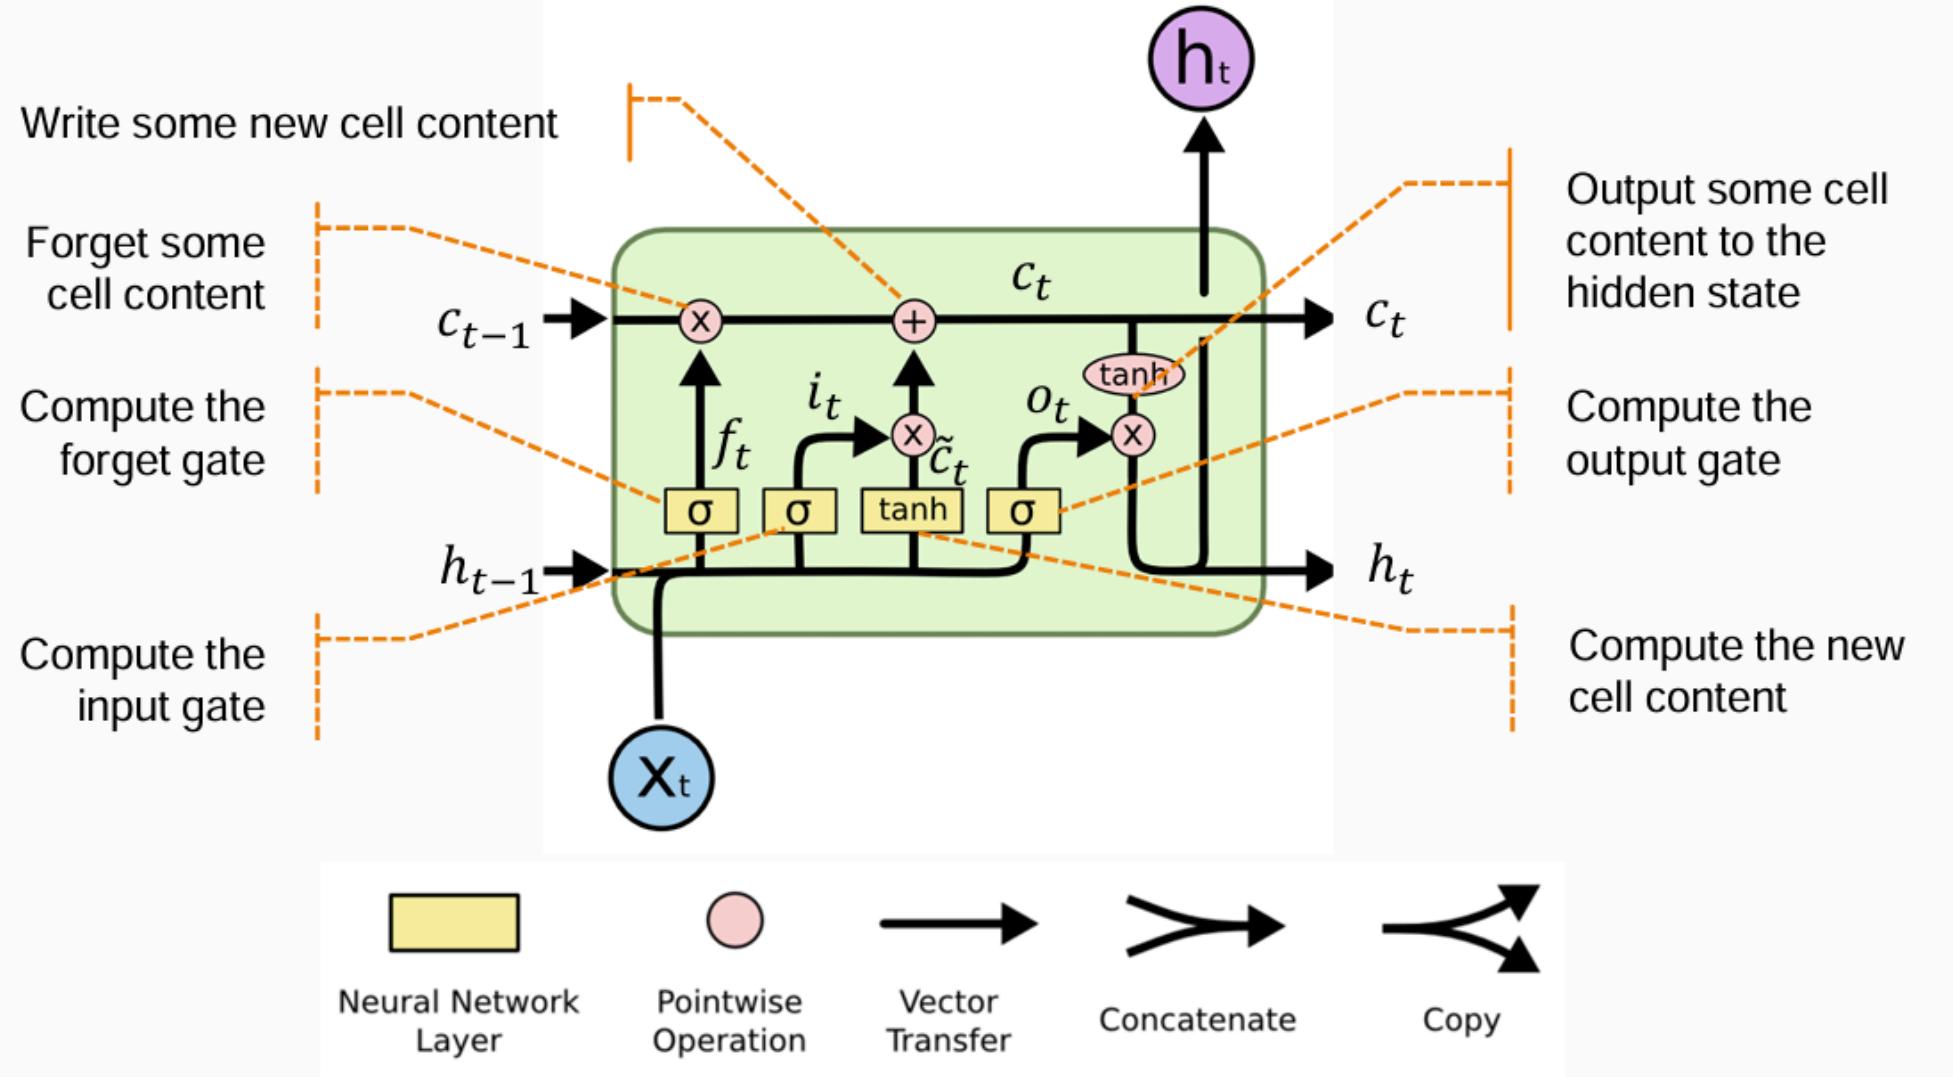
Vector  
Transfer



Concatenate



Copy



## Thinkings

- *In N-grams, the model will crashed because of missing information of global context.  
Can we make N very large to solve this problem?*