

CS190C Lec5

Build Transformer Decoder


Attention please: These slides delete and only delete code implementation of all modules.

Overview


- Macro System Architecture
- Code Implementation of Basic Computational Components
- Code Implementation of RoPE
- Code Implementation of FFN and Attention
- Assembly of Transformer Block and Final Model

PART1: Macro System Architecture

Transformer Overall Architecture Diagram




Modules of the whole Transformer



- Pass through a **Token Embedding** module to turn to dense vector (Input layer)
- Pass through several **Transformer Block** modules to absorb information in multiple rounds (Hidden layer)
- Normalize the number scale of final tensor
- Linear transform, calculate possible scores for generating of each word (Output layer)

Modules of each Transformer Block





- Pre-layer RMSNorm module
- MHA Module with Residual Connection
 - contains RoPE module if position embedding needed
 - contains Softmax module for calculation of attention score
- FFN Module with Residual Connection
 - contains SiLU module

Moreover, almost all modules require calling the **linear transformation** module!

PART2: Code Implementation of Basic Computational Components

Basic Components to be Implemented



Modules to be implemented:

- Token Embedding
- Linear
- RMSNorm
- SiLU
- Softmax

1. class Generate_EMBEDDINGS

IDEA:

- Most original input: BPE encoding results (e.g., $[3, 10, 2, 6, 4]$, all token IDs from the vocabulary)
- Tensor shape: $[batch_size, seq_len]$
- Expected model input: Different words in the vocabulary have different embedding vectors $\text{emb}_i \in R^{d_model}$
- Implementation idea:
 - Generate a learnable matrix $W_e \in R^{|V|*d_model}$.
 - The i -th column represents emb_i
 - Randomly initialized W_e , representing no prior knowledge about the meaning of any word at first.

1. `class Generate_EMBEDDINGS`

Parameter scheme:

- Initialization phase: Pass in `vocab_size`, `d_model`, `device` (the device where PyTorch tensors are stored), and `dtype` (numerical type of tensor elements).
- Forward phase:
 - Pass in `token_ids` (shape `[batch_size, seq_len]`)
 - Output shape is `[batch_size, seq_len, d_model]`

1. **class Generate_EMBEDDINGS**

Code here

2. class Linear_Transform

IDEA:

- Assume we need to transform a 3-dimensional tensor into a 6-dimensional one...
- Mathematically speaking: Let the 3D tensor x ($1 * 3$) right-multiply a $3 * 6$ matrix w
- The shape of xw is then $1 * 6$, just like the diagram below.




Question: Linear operations are the most frequent in LLMs.....

Can this operation be accelerated as much as possible?


2. class Linear_Transform

- PyTorch tensors follow the "last dimension elements have contiguous memory addresses" principle.
- For a $3 * 6$ tensor W, its memory layout as the diagram(same color means contiguous):



- For a $4 * 3 * 6$ tensor ([batch size, rows, columns]), every 6 elements are also contiguous (e.g., the memory address difference between `[1,1,1]` and `[1,1,2]` is 1 units, and the memory address difference between `[1,1,1]` and `[1,2,1]` is 6 units).

2. class Linear_Transform




Perform 6 vector dot product operations. In each operation:

- x participates as a whole row, memory is contiguous, can fully utilize cache
- w participates as a whole column, memory is not contiguous, may not utilize cache effectively

How can we make w also participate "as a whole row" in each operation?

2. class Linear_Transform



- PyTorch's transpose operation does not change the tensor's underlying memory space, just change stepsize. (Can you give an example?)
- Create new `w : [6, 3]` (every 3 elements are contiguous in memory)
- Transpose `w` to `[3, 6]` : Can perform matrix multiplication, and the memory distribution remains unchanged
- `x` and `w` have fully contiguous memory access during each operation, allowing full utilization of cache!

2. **class Linear_Transform**

Code here

3. class RMSNorm

IDEA: Normalize the input tensor \vec{a} (We've discussed why at Lec3)

- $a_i = \frac{a_i}{RMS(\vec{a})} g_i$ (divide by normalization weight uniformly, and apply learnable fine-tuning)
- g_i is a learnable parameter
- $RMS(\vec{a}) = \sqrt{\left(\frac{1}{d_{model}} \sum a_i^2\right) + \epsilon}$, that is L2-Norm.

Input tensor shape: [batch_size, seq_len, d_model] \Rightarrow Not a 1D vector, how to handle?


PyTorch's broadcasting mechanism:

- Operations are performed on the last few dimensions by default, previous dimensions are all replication operations

3. class RMSNorm

Code here

4. class SiLU_Activation



$$\text{SiLU}: f(x) = x \cdot \sigma(x)$$

$$\text{ReLU}: f(x) = \max(0, x)$$

- For $x < 0$: roughly equals to 0
- For $x > 0$: roughly stays the same
- Compare with ReLU: smooth and differentiable at $x = 0$

4. class SiLU_Activation

Code here

5. class Softmax_Activation

How to turn a score tensor to a distribution?

$$x_i = \frac{e^{x_i}}{\sum e^{x_i}}$$

- Each x_i calculates its exponential as a weight, then performs weight normalization
- Can make the advantage of relatively larger values more pronounced
- Even smaller values remain non-zero after normalization

Problem: What if there exists a very large x_i ? (e.g., normalizing [20, 3, 1005])

- $e^{1000} = \text{NAN}$

5. class Softmax_Activation

- Normalizing $[100, 101, 102]$ vs Normalizing $[-2, -1, 0]$
- Weight of 102: $\frac{e^{102}}{e^{102} + e^{101} + e^{100}} = \frac{e^0}{e^0 + e^{-1} + e^{-2}}$
- The Softmax normalization result of $[100, 101, 102]$ is equivalent to that of $[-2, -1, 0]$

Let x_{max} be the maximum value among x_i :

$$\begin{aligned}\text{Softmax}(x_i) &= \frac{e^{x_i}}{\sum e^{x_i}} \\ &= \frac{e^{x_i}/e^{x_{max}}}{\sum e^{x_i}/e^{x_{max}}} \\ &= \frac{e^{x_i-x_{max}}}{\sum e^{x_i-x_{max}}}\end{aligned}$$

That is: subtract x_{max} from all x_i to avoid problems with extremely large values that cannot be calculated!

5. **class Softmax_Activation**

Code here

PART3: Code Implementation of RoPE

Review: RoPE Calculation Rules

$$\vec{x} \Rightarrow R_i \vec{x} \ (\vec{x} \in R^d, d \text{ is even})$$

Divide the d-dimensional vector into several sub-segments of length 2, resulting in a total of $d/2$ sub-segments, each sub-segment makes a rotation of angle $\theta_{i,k}$.
(Proportional to position i)

$$R^i = \begin{bmatrix} R_1^i & 0 & 0 & \cdots & 0 \\ 0 & R_2^i & 0 & \cdots & 0 \\ 0 & 0 & R_3^i & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & R_{d/2}^i \end{bmatrix} \quad R_k^i = \begin{bmatrix} \cos(\theta_{i,k}) & -\sin(\theta_{i,k}) \\ \sin(\theta_{i,k}) & \cos(\theta_{i,k}) \end{bmatrix}.$$

$$\text{where } \theta_{i,k} = \frac{i}{\Theta^{2k/d}}$$

Review: RoPE Calculation Rules

$$R^i = \begin{bmatrix} R_1^i & 0 & 0 & \cdots & 0 \\ 0 & R_2^i & 0 & \cdots & 0 \\ 0 & 0 & R_3^i & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & R_{d/2}^i \end{bmatrix} \quad R_k^i = \begin{bmatrix} \cos(\theta_{i,k}) & -\sin(\theta_{i,k}) \\ \sin(\theta_{i,k}) & \cos(\theta_{i,k}) \end{bmatrix}.$$

Property of matrix R : $(R^m)^T R^n = R^{n-m}$


In the attention mechanism, for the query vector q_i at position i and the key vector k_j at position j :

- $q'_i = R^i q_i, k'_j = R^j k_j$
- $q_i'^T k'_j = q_i^T (R^i)^T R^j k_j = q_i^T R^{j-i} k_j$


How to Avoid Brute-force Matrix Multiplication?

- Observation: Even rows of the matrix correspond to the same transformation rule:
[cos, -sin], with angle $\theta_{i,k}$ as the variable
- If these even rows sharing the same rule can be computed efficiently in a unified manner, how should subsequent processing proceed?

$$R^i = \begin{bmatrix} R_1^i & 0 & 0 & \cdots & 0 \\ 0 & R_2^i & 0 & \cdots & 0 \\ 0 & 0 & R_3^i & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & R_{d/2}^i \end{bmatrix}$$




$$R_k^i = \begin{bmatrix} \cos(\theta_{i,k}) & -\sin(\theta_{i,k}) \\ \sin(\theta_{i,k}) & \cos(\theta_{i,k}) \end{bmatrix}.$$



How to Avoid Brute-force Matrix Multiplication?


- Calculate products of Red-Red , Yellow-Yellow , Blue-Blue , which correspond to: x_0, x_2, x_4 after RoPE
- x_1, x_3, x_5 are the same.
- Calculate each 2-number blocks, and calculate products.



How to Calculate 2-number blocks of R_i ?


- Pre-calculate a $\theta_{i,k}$ lookup table of size [max_seq_len, d/2]
- Then calculate lookup tables of the same size for $\cos(\theta_{i,k})$ and $\sin(\theta_{i,k})$
- Obtain values at positions 1, 3, 5 directly from the cos lookup table
- Obtain values at positions 2, 4, 6 directly from the sin lookup table

cos	-sin	cos	-sin	cos	-sin
1	2				
		3	4		
				5	6



How to Calculate 2-number blocks of R_i ?

- Concatenate pairs 12, 34, 56 into 3 blocks, perform block-wise dot product with the 3 blocks of x , obtaining the values for the three even rows.
- Similarly, obtain the values for all odd rows of the transformed x , concatenate both to get the complete transformed x vector.



Code Implementation of RoPE

Code here

Code Implementation of RoPE

Code here

PART4: Code Implementation of FFN and Attention

1. class Feed_Forward_Network

Module algorithm (excluding residual connection):

- Input tensor x (d_{model} dimensional)
- Route 1: Transform via w_1 to d_{ff} dimensional, pass through $SiLU$ activation function to get new x
- Route 2: Transform via w_3 for another expansion, getting gated d_{ff} dimensional
- Element-wise multiplication of x and gated, getting new x
- Transform back to d_{model} dimensional via w_2

We call it $\text{SwiGLU} : \text{FFN}(x) = \text{SwiGLU}(x, W_1, W_2, W_3)$

1. **class Feed_Forward_Network**

Code here

1. **class Feed_Forward_Network**

Code here

2. `class Multihead_Attention`

Module algorithm (excluding residual connection):

- Input tensor `x` with size `[batch_size, seq_len, d_model]`
- Through three linear transformations, essentially linearly transforming each `d_model` dimensional word vector of `x` to get:
 - `Q`, `K` matrices `[batch_size, seq_len, num_heads*d_k]`
 - `V` matrix `[batch_size, seq_len, num_heads*d_v]`
- Perform `RoPE` positional encoding on `Q`, `K` matrices
- Generate attention upper triangular mask (to prevent breaking the "autoregressive" assumption)
- Use QKV matrices and attention mask to compute attention output

2. class Multihead_Attention

Method for calculating attention output:



- Multiply Q , K matrices to calculate token feature matching scores
- Apply mask processing to the matching score matrix
- $\text{Softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)$, calculate attention scores between tokens
- Multiply the score matrix by the v matrix to get the attention output of each head
- Multiply w_0 matrix to integrate heads

Calculation of Attention Output for One Head

Code here

Generation of Attention Mask

Code here



List of Sub-modules in Multihead_Attention

- Attention mask generation module
- Attention output calculation module (sdpa)
- RoPE module \Rightarrow Requires additional parameters like max_seq_len, theta, token_positions, etc.
- Four types of linear transformations: Q, K, V, O

Assembly of Complete Module

Code here

Assembly of Complete Module


Code here

Assembly of Complete Module

Code here

PART5: Assembly of Transformer Block and Final Model

Structure of Transformer Block



- A RMSNorm module
- A MHA module
- Residual connection
- An other RMSNorm module
- A FFN module
- Residual connection

The parameters received by each Block are the union of all parameters required by the above modules!


Assembly of Transformer Block

Code here

Assembly of Transformer Block

Code here

Structure of Full Transformer



- A Token Embedding module
- Several Transformer Block modules
- A RMSNorm module
- Final Linear Transformation module

Assembly of Complete Transformer

Code here

Code here

Assembly of Complete Transformer

Code here